

• NOTE

Click [here](#) to download the full example code

AUTOGRAD: AUTOMATIC DIFFERENTIATION

Central to all neural networks in PyTorch is the `autograd` package. Let's first briefly visit this, and we will then go to training our first neural network.

The `autograd` package provides automatic differentiation for all operations on Tensors. It is a define-by-run framework, which means that your backprop is defined by how your code is run, and that every single iteration can be different.

Let us see this in more simple terms with some examples.

TENSOR

`torch.Tensor` is the central class of the package. If you set its attribute `.requires_grad` as `True`, it starts to track all operations on it. When you finish your computation you can call `.backward()` and have all the gradients computed automatically. The gradient for this tensor will be accumulated into `.grad` attribute.

To stop a tensor from tracking history, you can call `.detach()` to detach it from the computation history, and to prevent future computation from being tracked.

To prevent tracking history (and using memory), you can also wrap the code block in `with torch.no_grad():`. This can be particularly helpful when evaluating a model because the model may have trainable parameters with `requires_grad=True`, but for which we don't need the gradients.

There's one more class which is very important for autograd implementation - a `Function`.

`Tensor` and `Function` are interconnected and build up an acyclic graph, that encodes a complete history of computation. Each tensor has a `.grad_fn` attribute that references a `Function` that has created the `Tensor` (except for Tensors created by the user - their `grad_fn` is `None`).

If you want to compute the derivatives, you can call `.backward()` on a `Tensor`. If `Tensor` is a scalar (i.e. it holds a one element data), you don't need to specify any arguments to `backward()`, however if it has more elements, you need to specify a `gradient` argument that is a tensor of matching shape.

```
import torch
```

Create a tensor and set `requires_grad=True` to track computation with it

```
x = torch.ones(2, 2, requires_grad=True)
print(x)
```

Out:

```
tensor([[1., 1.],
        [1., 1.]], requires_grad=True)
```

Do an operation of tensor:

```
y = x + 2
print(y)
```

Out:

```
tensor([[3., 3.],
        [3., 3.]], grad_fn=<AddBackward0>)
```

`y` was created as a result of an operation, so it has a `grad_fn`.

```
print(y.grad_fn)
```

Out:

```
<AddBackward0 object at 0x7f1b248453c8>
```

Do more operations on y

```
z = y * y * 3
out = z.mean()

print(z, out)
```

Out:

```
tensor([[27., 27.],
        [27., 27.]], grad_fn=<MulBackward0>) tensor(27., grad_fn=<MeanBackward0>)
```

`.requires_grad_(...)` changes an existing Tensor’s `requires_grad` flag in-place. The input flag defaults to `False` if not given.

```
a = torch.randn(2, 2)
a = ((a * 3) / (a - 1))
print(a.requires_grad)
a.requires_grad_()
print(a.requires_grad)
b = (a * a).sum()
print(b.grad_fn)
```

Out:

```
False
True
<SumBackward0 object at 0x7f1b24845f98>
```

GRADIENTS

Let’s backprop now Because `out` contains a single scalar, `out.backward()` is equivalent to `out.backward(torch.tensor(1.))`.

```
out.backward()
```

print gradients d(out)/dx

```
print(x.grad)
```

Out:

```
tensor([[4.5000, 4.5000],
        [4.5000, 4.5000]])
```

You should have got a matrix of `4.5`. Let’s call the `out` Tensor “*o*”. We have that $o = \frac{1}{4} \sum_i z_i$, $z_i = 3(x_i + 2)^2$ and $z_i|_{x_i=1} = 27$. Therefore, $\frac{\partial o}{\partial x_i} = \frac{3}{2}(x_i + 2)$, hence $\frac{\partial o}{\partial x_i}|_{x_i=1} = \frac{9}{2} = 4.5$.

Mathematically, if you have a vector valued function $\vec{y} = f(\vec{x})$, then the gradient of \vec{y} with respect to \vec{x} is a Jacobian matrix:

$$J = \begin{pmatrix} \frac{\partial y_1}{\partial x_1} & \cdots & \frac{\partial y_m}{\partial x_1} \\ \vdots & \ddots & \vdots \\ \frac{\partial y_1}{\partial x_n} & \cdots & \frac{\partial y_m}{\partial x_n} \end{pmatrix}$$

Generally speaking, `torch.autograd` is an engine for computing Jacobian-vector product. That is, given any vector $v = (v_1 \ v_2 \ \cdots \ v_m)^T$, compute the product $J \cdot v$. If v happens to be the gradient of a scalar function $l = g(\vec{y})$, that is, $v = \left(\frac{\partial l}{\partial y_1} \ \cdots \ \frac{\partial l}{\partial y_m} \right)^T$, then by the chain rule, the Jacobian-vector product would be the gradient of l with respect to \vec{x} :

$$J \cdot v = \begin{pmatrix} \frac{\partial y_1}{\partial x_1} & \cdots & \frac{\partial y_m}{\partial x_1} \\ \vdots & \ddots & \vdots \\ \frac{\partial y_1}{\partial x_n} & \cdots & \frac{\partial y_m}{\partial x_n} \end{pmatrix} \begin{pmatrix} \frac{\partial l}{\partial y_1} \\ \vdots \\ \frac{\partial l}{\partial y_m} \end{pmatrix} = \begin{pmatrix} \frac{\partial l}{\partial x_1} \\ \vdots \\ \frac{\partial l}{\partial x_n} \end{pmatrix}$$

This characteristic of Jacobian-vector product makes it very convenient to feed external gradients into a model that has non-scalar output.

Now let’s take a look at an example of Jacobian-vector product:

```
x = torch.randn(3, requires_grad=True)

y = x * 2
while y.data.norm() < 1000:
    y = y * 2

print(y)
```

Out:

```
tensor([-278.6740,  935.4016,  439.6572], grad_fn=<MulBackward0>)
```

Now in this case `y` is no longer a scalar. `torch.autograd` could not compute the full Jacobian directly, but if we just want the Jacobian-vector product, simply pass the vector to `backward` as argument:

```
v = torch.tensor([0.1, 1.0, 0.0001], dtype=torch.float)
y.backward(v)

print(x.grad)
```

Out:

```
tensor([4.0960e+02, 4.0960e+03, 4.0960e-01])
```

You can also stop autograd from tracking history on Tensors with `.requires_grad=True` by wrapping the code block in `with torch.no_grad()` :

```
print(x.requires_grad)
print((x ** 2).requires_grad)

with torch.no_grad():
    print((x ** 2).requires_grad)
```

Out:

```
True
True
False
```

Read Later:

Documentation of `autograd` and `Function` is at <https://pytorch.org/docs/autograd>

Total running time of the script: (0 minutes 3.437 seconds)

⬇ **Download Python source code: autograd_tutorial.py**

⬇ **Download Jupyter notebook: autograd_tutorial.ipynb**

Gallery generated by Sphinx-Gallery

[< Previous](#)

[Next >](#)

© Copyright 2017, PyTorch.

Built with [Sphinx](#) using a [theme](#) provided by [Read the Docs](#).

Docs

Access comprehensive developer documentation for PyTorch

[View Docs >](#)

Tutorials

Get in-depth tutorials for beginners and advanced developers

[View Tutorials >](#)

Resources

Find development resources and get your questions answered

[View Resources >](#)



PyTorch

- Get Started
- Features
- Ecosystem
- Blog
- Resources

Support

- Tutorials
- Docs
- Discuss
- Github Issues
- Slack

Follow Us

[Email Address](#) [→](#)

